



Code

Profiling and

Optimization

Day 2

Memory layout and using
Debugging tools



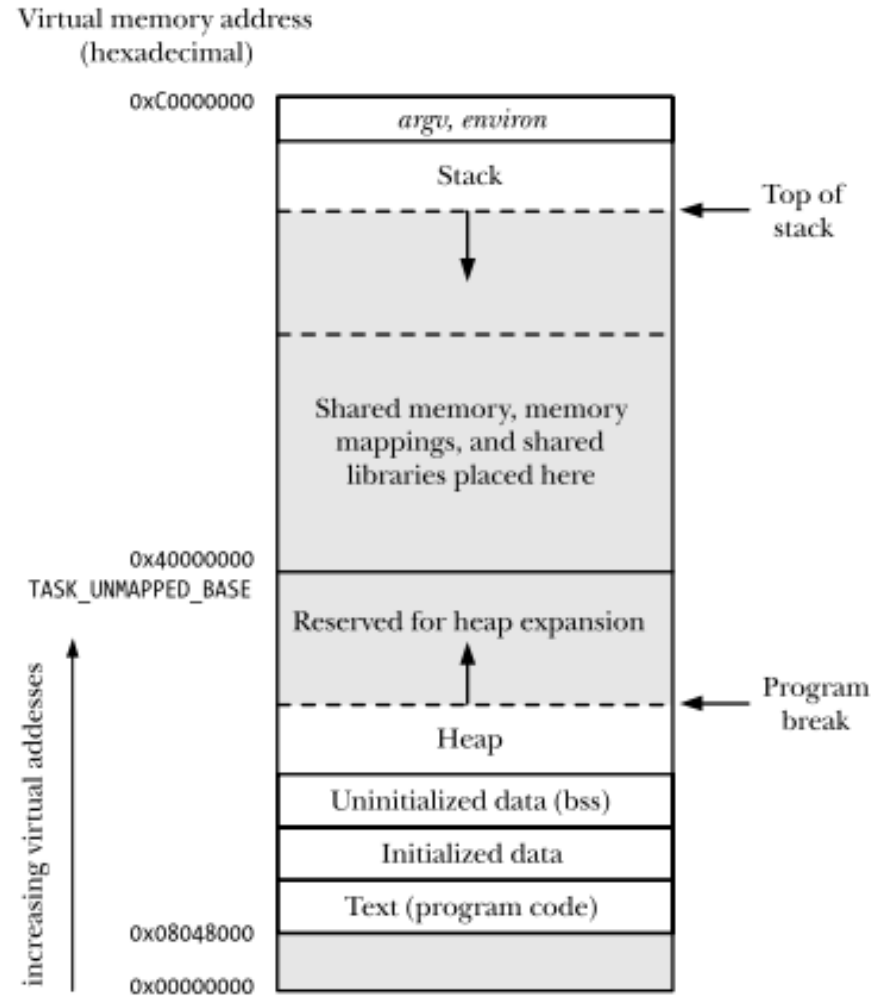


Figure 48-2: Locations of shared memory, memory mappings, and shared libraries (x86-32)

Memory Layout

Sections

- .text
- .data
- .bss
- .rodata
- Stack
- Heap



.data

- Global initialized variables
- For e.g.

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

// Global variables, initialized, will be allocated in the data segment
int a = 10;
int b = 20;
```



.data

```
#include <stdio.h>
#include <stdlib.h>
#include <time.h>

// Global variables, initialized, will be allocated in the data segment
int a = 10;
int b = 20;
```

- Global initialized variables
- For e.g.

```
_objdump -x -D -s a.out > a.dump
```

```
0000000000404034 g      0 .data 0000000000000004          a
```

```
0000000000404038 g      0 .data 0000000000000004          b
```

```
Contents of section .data:
```

```
404030 00000000 0a000000 14000000          .....
```



.bss

- Global uninitialized variables
- For e.g.

```
// Global variables, uninitialized, will be allocated in the bss segment  
int c;  
int d;
```



`.bss`

```
// Global variables, uninitialized, will be allocated in the bss segment  
int c;  
int d;
```

- Global uninitialized variables
- For e.g.

```
objdump -x -D -s a.out > a.dump
```

```
24 .bss 0000000c 000000000040403c 000000000040403c 0000303c 2**2
```

```
0000000000404040 g 0 .bss 0000000000000004 c
```

```
0000000000404044 g 0 .bss 0000000000000004 d
```



.rodata

- Read only data; Note if we try to modify .rodata, it will generate seg fault!
- For e.g.

```
✦ ✦ char *str = "Hello, World!";  
    printf("str = %s\n", str);  
}
```



.rodata

```
char *str = "Hello, World!";  
printf("str = %s\n", str);
```

- Read only data; Note if we try to modify .rodata, it will generate seg fault!
- For e.g.

Contents of section **.rodata**:

402000	01000200	00000000	00000000	00000000
402010	63203d20	25642c20	64203d20	25640a00	c = %d, d = %d..
402020	65203d20	25642c20	66203d20	25640a00	e = %d, f = %d..
402030	48656c6c	6f2c2057	6f726c64	21007374	Hello, World!.st
402040	72203d20	25730a00	61203d20	25642c20	r = %s..a = %d,
402050	62203d20	25640a00	456c656d	656e7473	b = %d..Elements
402060	206f6620	6172723a	20002564	2000	of arr: %d .



.text

- Code segment, instructions to be executed
- For e.g.

```
void foo(int p, int q)
{
    // initializing global variables
    c = p;
    d = q;
    printf("c = %d, d = %d\n", c, d);

    // defining local variables, will be allocated on the stack
    int e = 50;
    int f = 60;
    printf("e = %d, f = %d\n", e, f);

    char *str = "Hello, World!";
    printf("str = %s\n", str);
}
```

void foo
Click to co



.text

- Code segment, instructions to be executed
- For e.g.

```
void foo(int p, int q)
{
    // initializing global variables
    c = p;
    d = q;
    printf("c = %d, d = %d\n", c, d);
```

```
    // defining local variables, will be allocated on the stack
    int e = 50;
    int f = 60;
    printf("e = %d, f = %d\n", e, f);
```

```
    char *str = "Hello, World!";
    printf("str = %s\n", str);
}
```

void foo
Click to co

```
000000000401176 <foo>:
401176: 55                push    %rbp
401177: 48 89 e5          mov     %rsp,%rbp
40117a: 48 83 ec 20        sub     $0x20,%rsp
40117e: 89 7d ec          mov     %edi,-0x14(%rbp)
401181: 89 75 e8          mov     %esi,-0x18(%rbp)
401184: 8b 45 ec          mov     -0x14(%rbp),%eax
401187: 89 05 b3 2e 00 00 mov     %eax,0x2eb3(%rip)    # 404040 <c>
40118d: 8b 45 e8          mov     -0x18(%rbp),%eax
401190: 89 05 ae 2e 00 00 mov     %eax,0x2eae(%rip)    # 404044 <d>
401196: 8b 15 a8 2e 00 00 mov     0x2ea8(%rip),%edx    # 404044 <d>
40119c: 8b 05 9e 2e 00 00 mov     0x2e9e(%rip),%eax    # 404040 <c>
4011a2: 89 c6            mov     %eax,%esi
4011a4: bf 10 20 40 00    mov     $0x402010,%edi
4011a9: b8 00 00 00 00    mov     $0x0,%eax
4011ae: e8 8d fe ff ff    call    401040 <printf@plt>
4011b3: c7 45 fc 32 00 00 00 movl    $0x32,-0x4(%rbp)
4011ba: c7 45 f8 3c 00 00 00 movl    $0x3c,-0x8(%rbp)
4011c1: 8b 55 f8          mov     -0x8(%rbp),%edx
4011c4: 8b 45 fc          mov     -0x4(%rbp),%eax
4011c7: 89 c6            mov     %eax,%esi
4011c9: bf 20 20 40 00    mov     $0x402020,%edi
4011ce: b8 00 00 00 00    mov     $0x0,%eax
4011d3: e8 68 fe ff ff    call    401040 <printf@plt>
4011d8: 48 c7 45 f0 30 20 40 movq    $0x402030,-0x10(%rbp)
4011df: 00
4011e0: 48 8b 45 f0        mov     -0x10(%rbp),%rax
4011e4: 48 89 c6          mov     %rax,%rsi
4011e7: bf 3e 20 40 00    mov     $0x40203e,%edi
4011ec: b8 00 00 00 00    mov     $0x0,%eax
4011f1: e8 4a fe ff ff    call    401040 <printf@plt>
4011f6: 90                nop
4011f7: c9                leave
4011f8: c3                ret
```

stack

- Local variables will be adjusted in stack
- For e.g.

```
void foo(int p, int q)
{
    // initializing global variables
    c = p;
    d = q;
    printf("c = %d, d = %d\n", c, d);
```

```
    // defining local variables, will be allocated on the stack
    int e = 50;
    int f = 60;
    printf("e = %d, f = %d\n", e, f);
```

```
    char *str = "Hello, World!";
    printf("str = %s\n", str);
}
```

void foo
Click to co

```
000000000401176 <foo>:
401176: 55                push    %rbp
401177: 48 89 e5          mov     %rsp,%rbp
40117a: 48 83 ec 20       sub     $0x20,%rsp
40117e: 89 7d ec          mov     %edi,-0x14(%rbp)
401181: 89 75 e8          mov     %esi,-0x18(%rbp)
401184: 8b 45 ec          mov     -0x14(%rbp),%eax
401187: 89 05 b3 2e 00 00 mov     %eax,0x2eb3(%rip)    # 404040 <c>
40118d: 8b 45 e8          mov     -0x18(%rbp),%eax
401190: 89 05 ae 2e 00 00 mov     %eax,0x2eae(%rip)    # 404044 <d>
401196: 8b 15 a8 2e 00 00 mov     0x2ea8(%rip),%edx    # 404044 <d>
40119c: 8b 05 9e 2e 00 00 mov     0x2e9e(%rip),%eax    # 404040 <c>
4011a2: 89 c6            mov     %eax,%esi
4011a4: bf 10 20 40 00    mov     $0x402010,%edi
4011a9: b8 00 00 00 00    mov     $0x0,%eax
4011ae: e8 8d fe ff ff    call    401040 <printf@plt>
4011b3: c7 45 fc 32 00 00 00 movl    $0x32,-0x4(%rbp)
4011ba: c7 45 f8 3c 00 00 00 movl    $0x3c,-0x8(%rbp)
4011c1: 8b 55 f8          mov     -0x8(%rbp),%edx
4011c4: 8b 45 fc          mov     -0x4(%rbp),%eax
4011c7: 89 c6            mov     %eax,%esi
4011c9: bf 20 20 40 00    mov     $0x402020,%edi
4011ce: b8 00 00 00 00    mov     $0x0,%eax
4011d3: e8 68 fe ff ff    call    401040 <printf@plt>
4011d8: 48 c7 45 f0 30 20 40 movq    $0x402030,-0x10(%rbp)
4011df: 00
4011e0: 48 8b 45 f0       mov     -0x10(%rbp),%rax
4011e4: 48 89 c6          mov     %rax,%rsi
4011e7: bf 3e 20 40 00    mov     $0x40203e,%edi
4011ec: b8 00 00 00 00    mov     $0x0,%eax
4011f1: e8 4a fe ff ff    call    401040 <printf@plt>
4011f6: 90              nop
4011f7: c9              leave
4011f8: c3              ret
```

Allocates 32 bytes on the stack for local variables.

stack

- Local variables will be adjusted in stack
- For e.g.

```
void foo(int p, int q)
{
    // initializing global variables
    c = p;
    d = q;
    printf("c = %d, d = %d\n", c, d);

    // defining local variables, will be allocated on the stack
    int e = 50;
    int f = 60;
    printf("e = %d, f = %d\n", e, f);

    char *str = "Hello, World!";
    printf("str = %s\n", str);
}
```

void foo
Click to co

```
000000000401176 <foo>:
401176: 55                push    %rbp
401177: 48 89 e5          mov     %rsp,%rbp
40117a: 48 83 ec 20        sub     $0x20,%rsp
40117e: 89 7d ec          mov     %edi,-0x14(%rbp)
401181: 89 75 e8          mov     %esi,-0x18(%rbp)
401184: 8b 45 ec          mov     -0x14(%rbp),%eax
401187: 89 05 b3 2e 00 00 mov     %eax,0x2eb3(%rip)    # 404040 <c>
40118d: 8b 45 e8          mov     -0x18(%rbp),%eax
401190: 89 05 ae 2e 00 00 mov     %eax,0x2eae(%rip)    # 404044 <d>
401196: 8b 15 a8 2e 00 00 mov     0x2ea8(%rip),%edx    # 404044 <d>
40119c: 8b 05 9e 2e 00 00 mov     0x2e9e(%rip),%eax    # 404040 <c>
4011a2: 89 c6            mov     %eax,%esi
4011a4: bf 10 20 40 00    mov     $0x402010,%edi
4011a9: b8 00 00 00 00    mov     $0x0,%eax
4011ae: e8 8d fe ff ff    call    401040 <printf@plt>
4011b3: c7 45 fc 32 00 00 00 movl    $0x32,-0x4(%rbp)
4011ba: c7 45 f8 3c 00 00 00 movl    $0x3c,-0x8(%rbp)
4011c1: 8b 55 f8          mov     -0x8(%rbp),%edx
4011c4: 8b 45 fc          mov     -0x4(%rbp),%eax
4011c7: 89 c6            mov     %eax,%esi
4011c9: bf 20 20 40 00    mov     $0x402020,%edi
4011ce: b8 00 00 00 00    mov     $0x0,%eax
4011d3: e8 68 fe ff ff    call    401040 <printf@plt>
4011d8: 48 c7 45 f0 30 20 40 movq    $0x402030,-0x10(%rbp)
4011df: 00
4011e0: 48 8b 45 f0        mov     -0x10(%rbp),%rax
4011e4: 48 89 c6          mov     %rax,%rsi
4011e7: bf 3e 20 40 00    mov     $0x40203e,%edi
4011ec: b8 00 00 00 00    mov     $0x0,%eax
4011f1: e8 4a fe ff ff    call    401040 <printf@plt>
4011f6: 90              nop
4011f7: c9              leave
4011f8: c3              ret
```

stack

Allocates 32 bytes on the stack for local variables.

- Local variables will be adjusted in stack
- For e.g.

Stores the values of p and q.

```
void foo(int p, int q)
{
    // initializing global variables
    c = p;
    d = q;
    printf("c = %d, d = %d\n", c, d);
```

```
    // defining local variables, will be allocated on the stack
    int e = 50;
    int f = 60;
    printf("e = %d, f = %d\n", e, f);
```

```
    char *str = "Hello, World!";
    printf("str = %s\n", str);
}
```

void foo
Click to co

```
000000000401176 <foo>:
401176: 55                push    %rbp
401177: 48 89 e5          mov     %rsp,%rbp
40117a: 48 83 ec 20       sub     $0x20,%rsp
40117e: 89 7d ec          mov     %edi,-0x14(%rbp)
401181: 89 75 e8          mov     %esi,-0x18(%rbp)
401184: 8b 45 ec          mov     -0x14(%rbp),%eax
401187: 89 05 b3 2e 00 00 mov     %eax,0x2eb3(%rip)    # 404040 <c>
40118d: 8b 45 e8          mov     -0x18(%rbp),%eax
401190: 89 05 ae 2e 00 00 mov     %eax,0x2eae(%rip)    # 404044 <d>
401196: 8b 15 a8 2e 00 00 mov     0x2ea8(%rip),%edx    # 404044 <d>
40119c: 8b 05 9e 2e 00 00 mov     0x2e9e(%rip),%eax    # 404040 <c>
4011a2: 89 c6             mov     %eax,%esi
4011a4: bf 10 20 40 00    mov     $0x402010,%edi
4011a9: b8 00 00 00 00    mov     $0x0,%eax
4011ae: e8 8d fe ff ff    call    401040 <printf@plt>
4011b3: c7 45 fc 32 00 00 00 movl    $0x32,-0x4(%rbp)
4011ba: c7 45 f8 3c 00 00 00 movl    $0x3c,-0x8(%rbp)
4011c1: 8b 55 f8          mov     -0x8(%rbp),%edx
4011c4: 8b 45 fc          mov     -0x4(%rbp),%eax
4011c7: 89 c6             mov     %eax,%esi
4011c9: bf 20 20 40 00    mov     $0x402020,%edi
4011ce: b8 00 00 00 00    mov     $0x0,%eax
4011d3: e8 68 fe ff ff    call    401040 <printf@plt>
4011d8: 48 c7 45 f0 30 20 40 movq    $0x402030,-0x10(%rbp)
4011df: 00
4011e0: 48 8b 45 f0       mov     -0x10(%rbp),%rax
4011e4: 48 89 c6          mov     %rax,%rsi
4011e7: bf 3e 20 40 00    mov     $0x40203e,%edi
4011ec: b8 00 00 00 00    mov     $0x0,%eax
4011f1: e8 4a fe ff ff    call    401040 <printf@plt>
4011f6: 90               nop
4011f7: c9               leave
4011f8: c3               ret
```

heap

- Dynamically allocated variables will be in the heap
- For e.g.

```
void baz()
{
    // array of size 10 using malloc, so will be allocated on the heap
    int *arr = (int *)malloc(10 * sizeof(int));

    srand(time(NULL)); // seed the random number generator

    // Initialize the array randomly
    for (int i = 0; i < 10; i++)
    {
        arr[i] = rand() % 10;
    }
    printf("Elements of arr: ");
    for (int i = 0; i < 10; i++)
    {
        printf("%d ", arr[i]);
    }
    printf("\n");
}
```



heap

- Dynamically allocated variables will be in the heap
- For e.g.

```
void baz()
{
    // array of size 10 using malloc, so will be allocated on the heap
    int *arr = (int *)malloc(10 * sizeof(int));

    srand(time(NULL)); // seed the random number generator

    // Initialize the array randomly
    for (int i = 0; i < 10; i++)
    {
        arr[i] = rand() % 10;
    }

    printf("Elements of arr: ");
    for (int i = 0; i < 10; i++)
    {
        printf("%d ", arr[i]);
    }
    printf("\n");
}
```

Call to malloc@plt

```
000000000040121d <baz>:
40121d: 55                    push    %rbp
40121e: 48 89 e5             mov     %rsp,%rbp
401221: 48 83 ec 10          sub     $0x10,%rsp
401225: bf 28 00 00 00       mov     $0x28,%edi
40122a: e8 41 fe ff ff       call    401070 <malloc@plt>
40122f: 48 89 45 f0          mov     %rax,-0x10(%rbp)
401233: bf 00 00 00 00       mov     $0x0,%edi
401238: e8 23 fe ff ff       call    401060 <time@plt>
40123d: 89 c7                mov     %eax,%edi
40123f: e8 0c fe ff ff       call    401050 <srand@plt>
401244: c7 45 fc 00 00 00 00 movl    $0x0,-0x4(%rbp)
40124b: eb 49                jmp     401296 <baz+0x79>
40124d: e8 2e fe ff ff       call    401080 <rand@plt>
401252: 89 c1                mov     %eax,%ecx
401254: 8b 45 fc             mov     -0x4(%rbp),%eax
401257: 48 98                cltq
401259: 48 8d 14 85 00 00 00 lea     0x0(,%rax,4),%rdx
401260: 00
401261: 48 8b 45 f0          mov     -0x10(%rbp),%rax
401265: 48 8d 34 02          lea     (%rdx,%rax,1),%rsi
401269: 48 63 c1             movslq  %ecx,%rax
40126c: 48 69 c0 67 66 66 66 imul    $0x66666667,%rax,%rax
```


heap

- Dynamically allocated variables will be in the heap
- For e.g.

```
void baz()
{
    // array of size 10 using malloc, so will be allocated on the heap
    int *arr = (int *)malloc(10 * sizeof(int));

    srand(time(NULL)); // seed the random number generator

    // Initialize the array randomly
    for (int i = 0; i < 10; i++)
    {
        arr[i] = rand() % 10;
    }
    printf("Elements of arr: ");
    for (int i = 0; i < 10; i++)
    {
        printf("%d ", arr[i]);
    }
    printf("\n");
}
```

Call to malloc@plt

Jump to GLIBC

```
000000000040121d <baz>:
40121d: 55                push    %rbp
40121e: 48 89 e5          mov     %rsp,%rbp
401221: 48 83 ec 10       sub     $0x10,%rsp
401225: bf 28 00 00 00    mov     $0x28,%edi
40122a: e8 41 fe ff ff    call   401070 <malloc@plt>
40122f: 48 89 45 f0       mov     %rax,-0x10(%rbp)
401233: bf 00 00 00 00    mov     $0x0,%edi
401238: e8 23 fe ff ff    call   401060 <time@plt>
40123d: 89 c7            mov     %eax,%edi
40123f: e8 0c fe ff ff    call   401050 <srand@plt>
401244: c7 45 fc 00 00 00 00 movl    $0x0,-0x4(%rbp)
40124b: eb 49            jmp     401296 <baz+0x79>
40124d: e8 2e fe ff ff    call   401080 <rand@plt>
401252: 89 c1            mov     %eax,%ecx
401254: 8b 45 fc         mov     -0x4(%rbp),%eax
401257: 48 98            cltq
401259: 48 8d 14 85 00 00 00 lea     0x0(,%rax,4),%rdx
401260: 00
401261: 48 8b 45 f0       mov     -0x10(%rbp),%rax
401265: 48 8d 34 02       lea     (%rdx,%rax,1),%rsi
401269: 48 63 c1         movslq  %ecx,%rax
40126c: 48 69 c0 67 66 66 66 imul    $0x66666667,%rax,%rax
```

```
0000000000401070 <malloc@plt>:
401070: ff 25 aa 2f 00 00 jmp     *0x2faa(%rip)          # 404020 <malloc@GLIBC_2.2.5>
401076: 68 04 00 00 00    push    $0x4
40107b: e9 a0 ff ff ff    jmp     401020 <_init+0x20>
```

Virtual memory address
(hexadecimal)

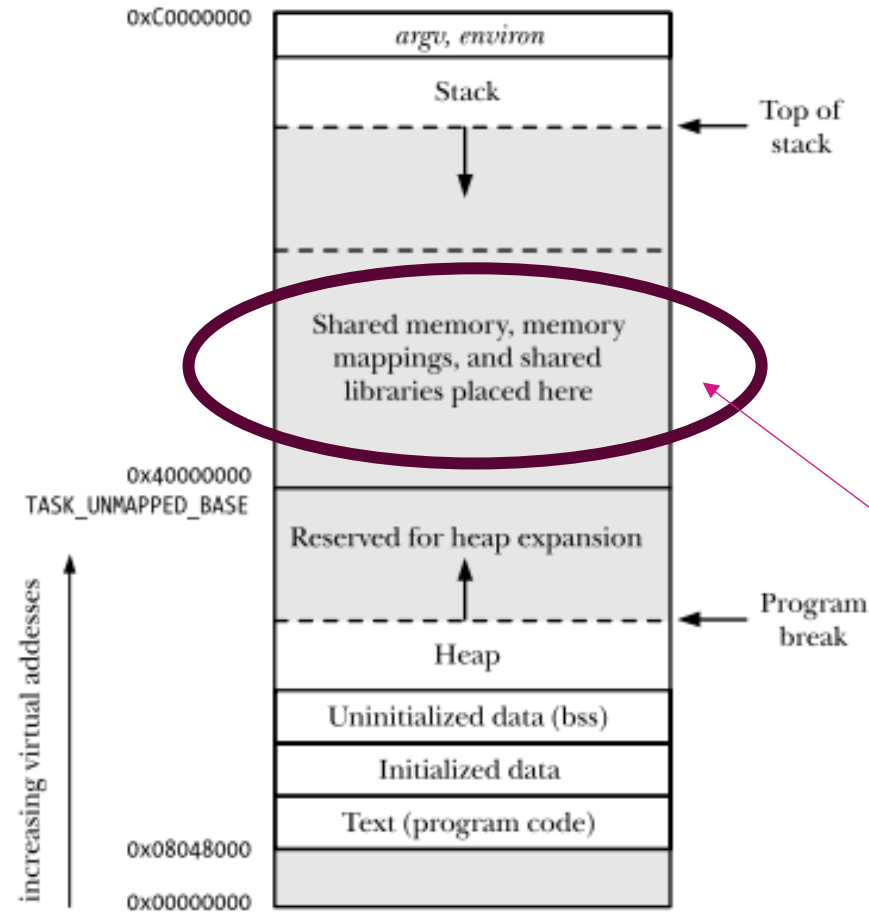


Figure 48-2: Locations of shared memory, memory mappings, and shared libraries (x86-32)

malloc@plt
be in the heap

To be handled by linker and loader

Jump to GLIBC

```
000000000040121d <baz>:
40121d: 55                push    %rbp
40121e: 48 89 e5          mov     %rsp,%rbp
401221: 48 83 ec 10       sub     $0x10,%rsp
401225: bf 28 00 00 00    mov     $0x28,%edi
40122a: e8 41 fe ff ff    call    401070 <malloc@plt>
40122f: 48 89 45 f0       mov     %rax,-0x10(%rbp)
401233: bf 00 00 00 00    mov     $0x0,%edi
401238: e8 23 fe ff ff    call    401060 <time@plt>
40123d: 89 c7             mov     %eax,%edi
40123f: e8 0c fe ff ff    call    401050 <srand@plt>
401244: c7 45 fc 00 00 00 00 movl    $0x0,-0x4(%rbp)
40124b: eb 49             jmp     401296 <baz+0x79>
40124d: e8 2e fe ff ff    call    401080 <rand@plt>
401252: 89 c1             mov     %eax,%ecx
401254: 8b 45 fc          mov     -0x4(%rbp),%eax
401257: 48 98             cltq
401259: 48 8d 14 85 00 00 00 lea     0x0(,%rax,4),%rdx
401260: 00
401261: 48 8b 45 f0       mov     -0x10(%rbp),%rax
401265: 48 8d 34 02       lea     (%rdx,%rax,1),%rsi
401269: 48 63 c1          movslq  %ecx,%rax
40126c: 48 69 c0 67 66 66 66 imul    $0x66666667,%rax,%rax
```

```
0000000000401070 <malloc@plt>:
401070: ff 25 aa 2f 00 00    jmp     *0x2faa(%rip)          # 404020 <malloc@GLIBC_2.2.5>
401076: 68 04 00 00 00      push    $0x4
40107b: e9 a0 ff ff ff      jmp     401020 <_init+0x20>
```

