

Code Profiling

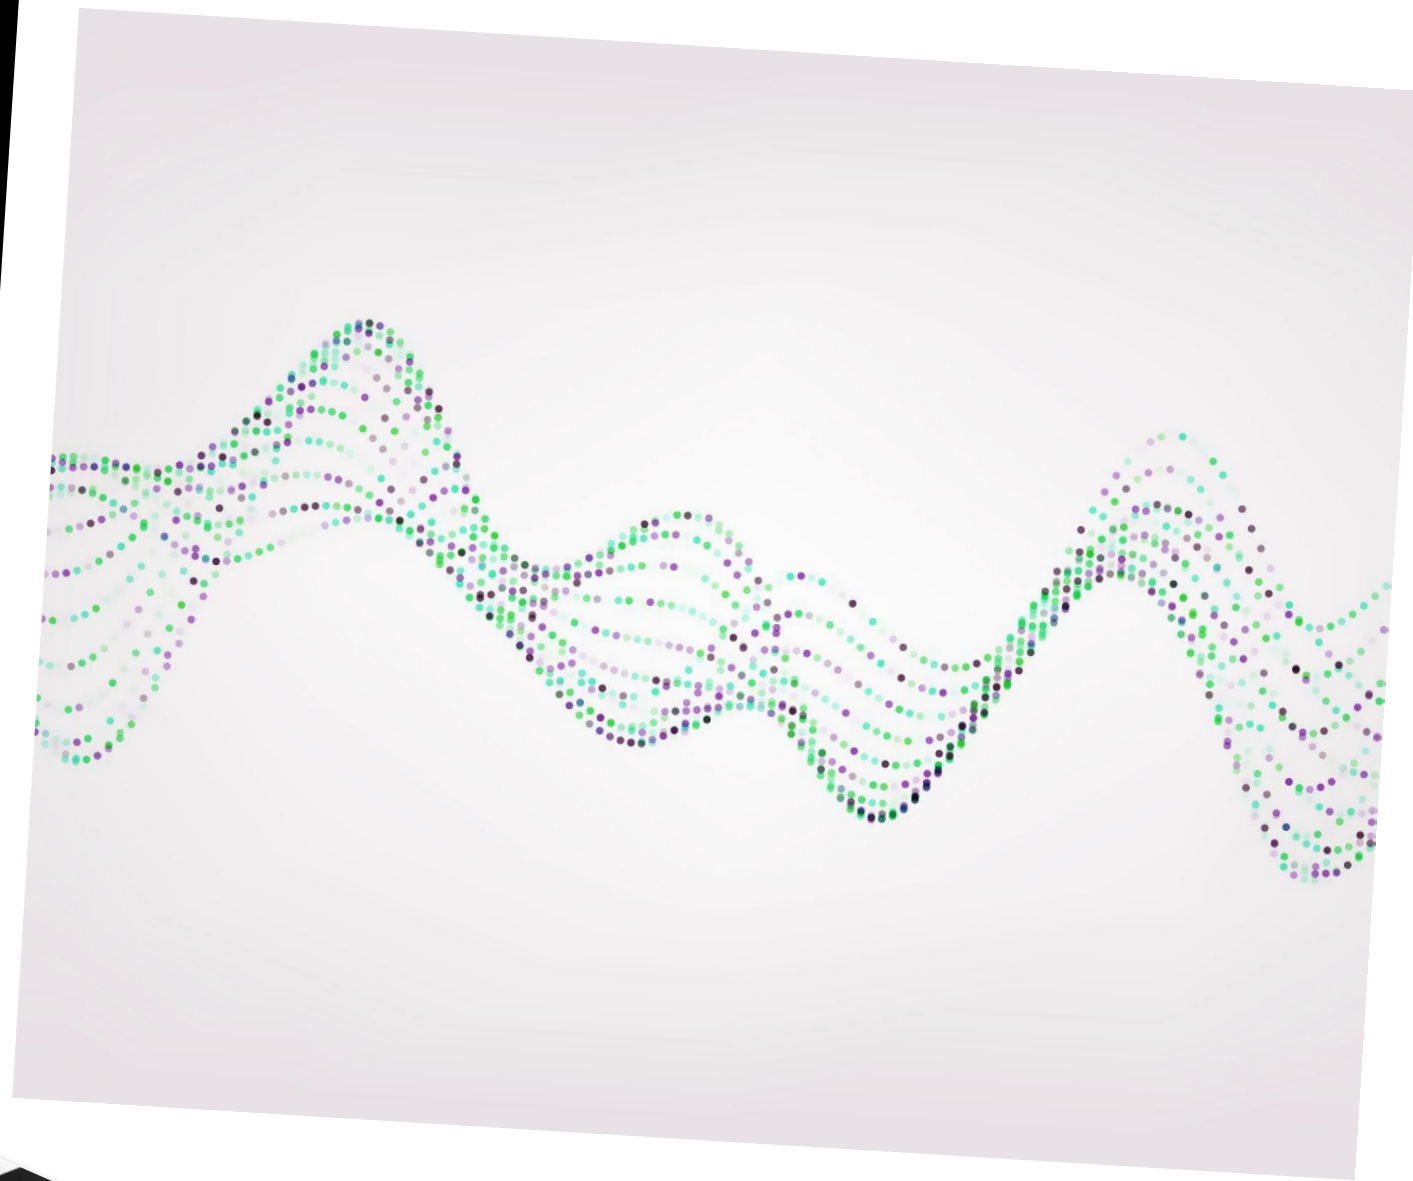
&

Optimization

Day 4

Measuring Performance Stats
with Linux Perf

Subhrajit & Pratyush



Getting started with Linux Perf

- a lightweight command-line utility for profiling and monitoring CPU performance on Linux systems.
- The perf command, by default, requires sudo privileges.
- `perf <options> subcommand <options/arguments>`



Subcommand	Description
<code>annotate</code>	Reads <i>perf.data</i> and shows annotated code.
<code>list</code>	Lists all measurable events.
<code>stat</code>	Gathers performance statistics.
<code>record</code>	Records samples into <i>perf.data</i> .
<code>report</code>	Reads <i>perf.data</i> and displays the profile.
<code>script</code>	Reads <i>perf.data</i> and displays trace output.
<code>top</code>	Profiling tool.

Getting started with Linux Perf

```
perf <options> subcommand <options/arguments>
```



When profiling a CPU with the `perf` command, the typical workflow is to use:

1. `perf list` to find events.
2. `perf stat` to count the events.
3. `perf record` to write events to a file.
4. `perf report` to browse the recorded file.
5. `perf script` to dump events after processing.

Getting started with Linux Perf

```
perf <options> subcommand <options/arguments>
```



1. `perf list` to find events.

```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ perf list
```

List of pre-defined events (to be used in `-e` or `-M`):

branch-instructions OR branches	[Hardware event]
branch-misses	[Hardware event]
bus-cycles	[Hardware event]
cache-misses	[Hardware event]
cache-references	[Hardware event]
cpu-cycles OR cycles	[Hardware event]
instructions	[Hardware event]
ref-cycles	[Hardware event]
alignment-faults	[Software event]

```
:|
```

Perf list

`perf <options> subcommand <options/arguments>`



2. `perf stat` to count the events.

To display CPU performance statistics for all standard CPU-wide hardware and software events, run:

```
● subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -a sleep 5

Performance counter stats for 'system wide':

      20,006.66 msec cpu-clock             #    4.000 CPUs utilized
           3,018      context-switches   #   150.850 /sec
            213      cpu-migrations      #    10.646 /sec
           5,690      page-faults        #   284.405 /sec
    3,26,48,27,377    cycles               #    0.163 GHz
    3,19,64,32,402    instructions        #    0.98  insn per cycle
    14,60,60,621     branches            #    7.301 M/sec
     20,23,904       branch-misses       #    1.39% of all branches
                    TopdownL1            #   58.4 % tma_backend_bound
                                         #    1.5 % tma_bad_speculation
                                         #   15.5 % tma_frontend_bound
                                         #   24.6 % tma_retiring

5.001729209 seconds time elapsed
```

Perf stat

`perf <options> subcommand <options/arguments>`



2. `perf stat` to count the events.

```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ perf stat ls
01_compilers 02_compilers 03_measurements 04_perf readme.md
```

Performance counter stats for 'ls':

0.64 msec	task-clock:u	#	0.631 CPUs utilized
0	context-switches:u	#	0.000 /sec
0	cpu-migrations:u	#	0.000 /sec
93	page-faults:u	#	144.314 K/sec
5,48,412	cycles:u	#	0.851 GHz
5,33,486	instructions:u	#	0.97 insn per cycle
1,05,345	branches:u	#	163.470 M/sec
6,247	branch-misses:u	#	5.93% of all branches
	TopdownL1	#	21.9 % tma_backend_bound
		#	23.0 % tma_bad_speculation
		#	35.5 % tma_frontend_bound
		#	19.6 % tma_retiring

0.001021751 seconds time elapsed

0.000000000 seconds user

0.001026000 seconds sys

Perf stat

`perf <options> subcommand <options/arguments>`



2. `perf stat` to count the events.

• `subhrajit@fedora:~/Code-Profiling-and-Optimization$ ps -a`

PID	TTY	TIME	CMD
71723	tty1	02:23:49	kwin_wayland
71730	tty1	05:20:46	sddm-greeter-qt
71767	tty1	00:00:00	maliit-keyboard
93487	pts/3	00:00:00	perf
93488	pts/3	00:00:00	less
94428	pts/4	00:00:00	top
94553	pts/3	00:00:00	ps

• `subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -p 94428 sleep 5`

Performance counter stats for process id '94428':

13.96 msec	task-clock	#	0.003 CPUs utilized
2	context-switches	#	143.297 /sec
0	cpu-migrations	#	0.000 /sec
0	page-faults	#	0.000 /sec
1,80,92,375	cycles	#	1.296 GHz
3,64,99,048	instructions	#	2.02 insn per cycle
75,09,874	branches	#	538.071 M/sec
49,175	branch-misses	#	0.65% of all branches
	TopdownL1	#	17.9 % tma_backend_bound
		#	5.6 % tma_bad_speculation
		#	36.8 % tma_frontend_bound
		#	39.6 % tma_retiring

5.001322737 seconds time elapsed

`perf <options> subcommand <options/arguments>`

• `subhrajit@fedora:~/Code-Profiling-and-Optimization$`

Perf stat


```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -e cycles -p 94428 sleep 5
```

Performance counter stats for process id '94428':

1,86,31,724 cycles

5.001649877 seconds time elapsed

```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -e cache-misses -p 94428 sleep 5
```

Performance counter stats for process id '94428':

1,37,889 cache-misses

5.001684361 seconds time elapsed

```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -e L1-dcache-loads ls  
01_compilers 02_compilers 03_measurements 04_perf readme.md
```

Performance counter stats for 'ls':

5,18,088 L1-dcache-loads

0.000973913 seconds time elapsed

0.001025000 seconds user

0.000000000 seconds sys

Perf stat

perf <options> subcommand <options/arguments>



```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -e cycles -p 94428 sleep 5
```

Performance counter stats for process id '94428':

1,86,31,724 cycles

5.001649877 seconds time elapsed

```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -e cache-misses -p 94428 sleep 5
```

Performance counter stats for process id '94428':

1,37,889 cache-misses

5.001684361 seconds time elapsed

```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -e L1-dcache-loads ls  
01_compilers 02_compilers 03_measurements 04_perf readme.md
```

Performance counter stats for 'ls':

5,18,088 L1-dcache-loads

0.000973913 seconds time elapsed

0.001025000 seconds user

0.000000000 seconds sys

```
subhrajit@fedora:~/Code-Profiling-and-Optimization$ sudo perf stat -e L1-dcache-load-misses ls  
01_compilers 02_compilers 03_measurements 04_perf readme.md
```

Performance counter stats for 'ls':

25,932 L1-dcache-load-misses

0.000904312 seconds time elapsed

0.000000000 seconds user

0.000961000 seconds sys



3. `perf record` to write events to a file.

- `subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf$ sudo perf record sleep 5`
[perf record: Woken up 1 times to write data]
[perf record: Captured and wrote 0.020 MB perf.data (8 samples)]
- `subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf$`

`subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf$ sudo perf report`

4. `perf report` to browse the recorded file.

```
Samples: 8 of event 'cycles:P', Event count (approx.): 1576135
Overhead Command Shared Object Symbol
 88.02% sleep [kernel.kallsyms] [k] _atomic_dec_and_lock
 11.21% sleep [kernel.kallsyms] [k] security_bprm_committing_creds
  0.72% perf-ex [kernel.kallsyms] [k] acpi_os_read_memory
  0.05% perf-ex [kernel.kallsyms] [k] native_write_msr
```

```
○ subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf$ sudo perf report --stdio
# To display the perf.data header info, please use --header/--header-only options.
#
#
# Total Lost Samples: 0
#
# Samples: 8 of event 'cycles:P'
# Event count (approx.): 1576135
#
# Overhead Command Shared Object Symbol
# .....
#
 88.02% sleep [kernel.kallsyms] [k] _atomic_dec_and_lock
 11.21% sleep [kernel.kallsyms] [k] security_bprm_committing_creds
  0.72% perf-ex [kernel.kallsyms] [k] acpi_os_read_memory
  0.05% perf-ex [kernel.kallsyms] [k] native_write_msr
#
# (Tip: To collect Processor Trace with samples use perf record -e '{intel_pt//,cycles}' ; perf script --call
#
```

Perf record

`perf <options> subcommand <options/arguments>`

5. `perf script` to dump events after processing.

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf$ sudo perf script
perf-exec 98596 220340.459016:      1 cycles:P: ffffffff830a9158 native_write_msr+0x8 ([kernel.>
perf-exec 98596 220340.459022:      1 cycles:P: ffffffff830a9158 native_write_msr+0x8 ([kernel.>
perf-exec 98596 220340.459031:      4 cycles:P: ffffffff830a9158 native_write_msr+0x8 ([kernel.>
perf-exec 98596 220340.459033:     50 cycles:P: ffffffff830a9158 native_write_msr+0x8 ([kernel.>
perf-exec 98596 220340.459035:     694 cycles:P: ffffffff830a915a native_write_msr+0xa ([kernel.>
perf-exec 98596 220340.459038:    11387 cycles:P: ffffffff839e8a98 acpi_os_read_memory+0x58 ([ker>
sleep 98596 220340.459043:    176714 cycles:P: ffffffff83750cca security_bprm_committing_creds>
sleep 98596 220340.459107:   1387284 cycles:P: ffffffff840c27f0 _atomic_dec_and_lock+0x0 ([ker>
(END)
```

The output prints the `perf.data` details in time order. Use the `script` subcommand as post-processing data.

Perf script

`perf <options> subcommand <options/arguments>`



Sample Code for Perf

- subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1\$ ls
sample.c
- subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1\$ gcc sample.c
- subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1\$ sudo perf stat ./a.out
x = 100000

Performance counter stats for './a.out':

0.43 msec	task-clock	#	0.568 CPUs utilized
0	context-switches	#	0.000 /sec
0	cpu-migrations	#	0.000 /sec
61	page-faults	#	142.147 K/sec
11,71,736	cycles	#	2.730 GHz
14,43,183	instructions	#	1.23 insn per cycle
2,92,781	branches	#	682.260 M/sec
6,232	branch-misses	#	2.13% of all branches
	TopdownL1	#	27.4 % tma_backend_bound
		#	8.8 % tma_bad_speculation
		#	35.6 % tma_frontend_bound
		#	28.2 % tma_retiring

0.000756136 seconds time elapsed

0.000000000 seconds user

0.000776000 seconds sys

```
#include <stdio.h>

int sample_function(int p)
{
    int n = 0;
    for (int i = 0; i < p; i++)
    {
        n = n + 1;
    }
    return n;
}

int main()
{
    int x = sample_function(100000);
    printf("x = %d\n", x);

    return 0;
}
```



Sample Code for Perf

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ gcc -g sample.c
```

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ sudo perf record -e cycles,instructions:u,instructions:k,LLC-loads,LLC-load-misses,branch-misses,bus-cycles,alignment-faults,page-faults ./a.out
```

```
x = 100000
```

```
[ perf record: Woken up 1 times to wr
```

```
[ perf record: Captured and wrote 0.0
```

```
#include <stdio.h>
```

```
int sample_function(int p)
```

```
{
    int n = 0;
    for (int i = 0; i < p; i++)
    {
        n = n + 1;
    }
    return n;
}
```

```
int main()
```

```
{
    int x = sample_function(100000);
    printf("x = %d\n", x);

    return 0;
}
```

```
Samples: 9 of event 'cycles', 4000 Hz, Event count (approx.): 2158827
```

```
sample function /home/subhrajit/Code-Profiling-and-Optimization/04_perf/demo1/a.out [Percent: local period]
```

Percent	int sample_function(int p)
	{
	push %rbp
	mov %rsp,%rbp
	mov %edi,-0x14(%rbp)
	int n = 0;
	movl \$0x0,-0x4(%rbp)
	for (int i = 0; i < p; i++)
	movl \$0x0,-0x8(%rbp)
	↓ jmp 1f
	{
	n = n + 1;
17:	addl \$0x1,-0x4(%rbp)
	for (int i = 0; i < p; i++)
	addl \$0x1,-0x8(%rbp)
1f:	mov -0x8(%rbp),%eax
	cmp -0x14(%rbp),%eax
	↑ jl 17
	}
	return n;
	mov -0x4(%rbp),%eax
	}
	pop %rbp
	← ret

100.00



Sample Code for Perf

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ gcc -g sample.c
```

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ sudo perf record -e cycles,instructions:u,instructions:k,LLC-loads,LLC-load-misses,branch-misses,bus-cycles,alignment-faults,  
x = 100000
```

```
[ perf record: Woken up 1 times
```

```
[ perf record: Captured and wrote
```

```
Samples: 10 of event 'instructions:u', 4000 Hz, Event count (approx.): 967587
```

```
sample function /home/subhrajit/Code-Profiling-and-Optimization/04_perf/demo1/a.out [Percent: local period]
```

```
#include <stdio.h>
```

```
int sample_function(int p)
```

```
{  
    int n = 0;  
    for (int i = 0; i < p; i++)  
    {  
        n = n + 1;  
    }  
    return n;  
}
```

```
int main()
```

```
{  
    int x = sample_function(100000);  
    printf("x = %d\n", x);  
  
    return 0;  
}
```

```
100.00
```

```
int sample_function(int p)
```

```
{  
    push %rbp  
    mov %rsp,%rbp  
    mov %edi,-0x14(%rbp)  
    int n = 0;  
    movl $0x0,-0x4(%rbp)  
    for (int i = 0; i < p; i++)  
        movl $0x0,-0x8(%rbp)  
        ↓ jmp 1f  
        {  
            n = n + 1;  
17:    addl $0x1,-0x4(%rbp)  
        }  
        for (int i = 0; i < p; i++)  
            addl $0x1,-0x8(%rbp)  
1f:    mov -0x8(%rbp),%eax  
        cmp -0x14(%rbp),%eax  
        ↑ jl 17  
    }  
    return n;  
    mov -0x4(%rbp),%eax  
}  
    pop %rbp  
← ret
```

Sample Code for Perf

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ gcc -g sample.c
```

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ sudo perf record -e cvcles.instructions:u.instructions:k.LLC-loads.LLC-load-misses.branch-misses,bus-cycles,alignment-faults,page-faults ./a.out
```

```
x = 100000
```

```
[ perf record: Woken up 1 times to write data ]
```

```
[ perf record: Captured and wrote 0.026 MB perf.data (65
```

```
#include <stdio.h>
```

```
int sample_function(int p)
```

```
{
    int n = 0;
    for (int i = 0; i < p; i++)
    {
        n = n + 1;
    }
    return n;
}
```

```
int main()
```

```
{
    int x = sample_function(100000);
    printf("x = %d\n", x);

    return 0;
}
```

```
Samples: 9 of event 'instructions:k', 4000 Hz, Event count (approx.): 1673815
refill obj stock /proc/kcore [Percent: local period]
```

Percent

Disassembly of section load0:

ffffffff834760f0 <load0>:

```
    nop
    push    %r15
    mov     %rdi,%r15
    push    %r14
    push    %r13
    push    %r12
    push    %rbp
    push    %rbx
    mov     %esi,%ebx
    pushf
    pop     %rax
    nop
    mov     %rax,%r12
    cli
    nop
    mov     %gs:0x7cba38fd(%rip),%rax
    and     $0x200,%r12d
    lea     0x325c0(%rax),%rbp
    mov     0x10(%rbp),%rax
    cmp     %rdi,%rax
    ↓ je     da
    mov     %rbp,%rdi
    → call   drain_obj_stock
```

100.00



Sample Code for Perf

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ gcc -g sample.c
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ sudo perf record -e cycles,instructions:u,instructions:k,LLC-loads,LLC-load-misses,branch-misses,bus-cycles,alignment-faults,page-faults ./a.out
x = 100000
[ perf record: Woken up 1 times to write data ]
[ perf record: Captured and wrote 0.026 MB perf.data ]
```

```
#include <stdio.h>

int sample_function(int p)
{
    int n = 0;
    for (int i = 0; i < p; i++)
    {
        n = n + 1;
    }
    return n;
}

int main()
{
    int x = sample_function(100000);
    printf("x = %d\n", x);

    return 0;
}
```

Samples: 8 of event 'LLC-loads', 4000 Hz, Event count (approx.): 4282
vma interval tree remove /proc/kcore [Percent: local period]

Percent	18d:	cmp %rdx,0x18(%rax)
	↑	jne 13c
	193:	test %rdi,%rdi
	↓	je 1b7
		mov %rbx,%rsi
		mov \$0xffffffff833c2f50,%rdx
		pop %rbx
		pop %rbp
		pop %r12
	→	jmp __rb_erase_color
	1ab:	mov %rcx,(%rbx)
		test %rcx,%rcx
	↓	jne 27e
	1b7:	pop %rbx
		pop %rbp
		pop %r12
	←	ret
		int3
		int3
		int3
		int3
	1c0:	testb \$0x1,(%rax)
		mov 0x10(%rax),%rcx
		mov %rdx,(%rax)



Sample Code for Perf

```
subhrajit@fedora:~/Code-Profiling-and-Optimization
subhrajit@fedora:~/Code-Profiling-and-Optimization
ses,bus-cycles,alignment-faults,page-faults ./a.out
x = 100000
[ perf record: Woken up 1 times to write data ]
[ perf record: Captured and wrote 0.026 MB perf.data ]
```

```
#include <stdio.h>

int sample_function(int p)
{
    int n = 0;
    for (int i = 0; i < p; i++)
    {
        n = n + 1;
    }
    return n;
}

int main()
{
    int x = sample_function(100000);
    printf("x = %d\n", x);

    return 0;
}
```

```
Samples: 7 of event 'LLC-load-misses', 4000 Hz, Event count (approx.): 1514
filemap_get_entry /proc/kcore [Percent: local period]
```

Percent		
	test	%rax,%rax
	↓ je	115
	test	\$0x1,%al
	↓ jne	115
	→ call	__rcu_read_lock
	nop	
	a3: mov	0x34(%rbx),%eax
	test	%eax,%eax
	↓ je	170
	mov	0x34(%rbx),%eax
	b1: → test	%eax,%eax
	↓ je	170
	lea	0x1(%rax),%edx
	lock	cmpxchg %edx,0x34(%rbx)
100.00	jne	b1
	→ call	__rcu_read_unlock
	mov	0x20(%rsp),%r13
	test	%r13,%r13
	↓ je	191
	movzbl	0x0(%r13),%ecx
	mov	0x10(%rsp),%rbp
	cmp	\$0x3f,%cl
	→ ja	filemap_get_entry.cold
	shr	%cl,%rbp
	and	\$0x3f,%ebp
	add	\$0x4,%rbp
	mov	0x8(%r13,%rbp,8),%rax
	mov	%rax,%rdx

LLC-load-misses,branch-mis



Sample Code for Perf

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ gcc -g sample.c
```

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04_perf/demo1$ sudo perf record -e cycles,instructions:u,instructions:k,LLC-loads,LLC-load-misses,branch-misses,bus-cycles,alignment-faults,page-faults ./a.out
```

```
x = 100000
```

```
[ perf record: Woken up 1 times to write data ]
```

```
[ perf record: Captured and wrote 0.026 MB perf.data ]
```

```
#include <stdio.h>
```

```
int sample_function(int p)
```

```
{
    int n = 0;
    for (int i = 0; i < p; i++)
    {
        n = n + 1;
    }
    return n;
}
```

```
int main()
```

```
{
    int x = sample_function(100000);
    printf("x = %d\n", x);

    return 0;
}
```

```
Samples: 8 of event 'branch-misses', 4000 Hz, Event count (approx.): 9069
```

```
mas next slot /proc/kcore [Percent: local period]
```

```
Percent
```

```
100.00  xor    %al,%al
        cmp    %rax,%r12
        ↓ je    406
        test   %rsi,%rsi
        ↓ jne   4b5
        test   %r14b,%r14b
        ↓ jne   4b3
        mov     0x10(%rbx),%rax
        cmp     %r15,%rax
        ↓ jae   4ac
        add     $0x1,%rax
        movzbl  0x3f(%rbx),%esi
        mov     %rax,0x8(%rbx)
        movzbl  0x3d(%rbx),%eax
        ↑ jmp    a4
124:     cmp     %r15,0x10(%rbx)
        ↓ jae   4ac
        mov     0x28(%rbx),%r8
        cmp     %r15,%r8
        ↓ jae   4ac
```



Sample Code for Perf

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04 perf/demo1$ gcc -g sample.c
```

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04 perf/demo1$ perf record -e d-misses,branch-misses
```

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/04 perf/demo1$ perf stat ./demo1
```

```
x = 100000
```

```
[ perf record: Woken up 1 times to write to the perf event
```

```
[ perf record: Captured and wrote 0.026 MB of perf data (100.00% of 262144 bytes)]
```

```
Samples: 7 of event 'bus-cycles', 4000 Hz, Event count (approx.): 14808
dl_relocate_object /usr/lib64/ld-linux-x86-64.so.2 [Percent: local period]
```

```
#include <stdio.h>
```

```
int sample_function(int p)
```

```
{
    int n = 0;
    for (int i = 0; i < p; i++)
    {
        n = n + 1;
    }
    return n;
}
```

```
int main()
```

```
{
    int x = sample_function(100000);
    printf("x = %d\n", x);

    return 0;
}
```

```
Percent 100.00 678:
          ↑ je      cd
          if (___glibc_unlikely (GLRO(dl_debug_mask) & DL_DEBUG_RELOC))
            test    %eax,%eax
          ↓ jne     25a0
          if (___glibc_unlikely (l->l_info[DT_TEXTREL] != NULL))
            mov     0xf0(%r14),%rax
            mov     %rax,-0xf8(%rbp)
            test    %rax,%rax
          ↓ jne     2564
            lazy = 0;
            xor     %r15d,%r15d
          ↑ jmp     13a
            nop
            ELF_DYNAMIC_RELOCATE (l, scope, lazy, consider_profiling, skip_ifunc);
          678: shr     $1,%rax
            add     $0x8,%rdx
          ↑ jmp     1d0
            nop
            + (((ElfW(Addr)) reloc_addr) - map->l_mach.gotplt) * 2;
          688: mov     %r12,%rdx
            sub     0x430(%r14),%rdx
            lea     (%rax,%rdx,2),%rax
            mov     %rax,(%r12)
```



Sample Code for Perf

```
subhrajit@fedora:~/Code-Profiling-and-Optimization/64-bit/ld-linux-x86-64.so.2
subhrajit@fedora:~/Code-Profiling-and-Optimization/64-bit/ld-linux-x86-64.so.2
ses,bus-cycles,alignment-faults,page-faults,cache-misses,branch-misses,load-misses,branch-misses
x = 100000
[ perf record: Woken up 1 times to write
[ perf record: Captured and wrote 0.026
```

```
#include <stdio.h>
```

```
int sample_function(int p)
```

```
{
    int n = 0;
    for (int i = 0; i < p; i++)
    {
        n = n + 1;
    }
    return n;
}
```

```
int main()
```

```
{
    int x = sample_function(100000);
    printf("x = %d\n", x);

    return 0;
}
```

```
Samples: 7 of event 'page-faults', 4000 Hz, Event count (approx.): 77
```

```
dl relocate object /usr/lib64/ld-linux-x86-64.so.2 [Percent: local period]
```

Percent

#endif

elf_machine_rel (map, scope, r, sym, rversion, r_addr_arg,

mov %r12,-0x78(%rbp)

if (__glibc_unlikely (r_type == R_X86_64_RELATIVE))

cmp \$0x8,%r15

↓ je f28

if (__glibc_unlikely (r_type == R_X86_64_RELATIVE64))

cmp \$0x26,%r15

↓ je f28

if (__glibc_unlikely (r_type == R_X86_64_NONE))

test %r15,%r15

↓ je 96b

if (ELFW(ST_BIND) ((*ref)->st_info) == STB_LOCAL

movzbl 0x4(%r12),%eax

mov %eax,%edi

shr \$0x4,%dil

↓ je ee0

movzbl 0x5(%r12),%edx

and \$0x3,%edx

sub \$0x1,%edx

|| __glibc_unlikely (dl_symbol_visibility_binds_local_p (*ref)))

cmp \$0x1,%edx

↓ jbe ee0

if (__glibc_unlikely (*ref == l->l_lookup_cache.sym)

cmp 0x440(%p12),%p12

C-load-misses,branch-misses

100.00



★ perf cheat sheet ★

JULIA EVANS
@b0rk

sourced from brendangregg.com/perf.html, which has many more great examples

important command line arguments

- a : entire system
- g : record stack traces
- e : choose an event to record
- p : specify a PID
- F : pick sample frequency

perf top: get updates live!

```
# Sample CPUs at 49 Hertz, show top symbols:
perf top -F 49

# Sample CPUs, show top process names and segments:
perf top -ns comm,dso

# Count system calls by process, refreshing every 1 second:
perf top -e raw_syscalls:sys_enter -ns comm -d 1

# Count sent network packets by process, rolling output:
stdbuf -oL perf top -e net:net_dev_xmit -ns comm | strings
```

sampling
tracing

perf stat: count events! CPU counters!

```
# CPU counter statistics for COMMAND:
perf stat COMMAND

# *Detailed* CPU counter statistics for COMMAND:
perf stat -ddd command

# Various basic CPU statistics, system wide:
perf stat -e cycles,instructions,cache-misses -a

# Count system calls for PID, until Ctrl-C:
perf stat -e 'syscalls:sys_enter_*' -p PID

# Count block device I/O events for the entire system, for 10 seconds:
perf stat -e 'block:*' -a sleep 10
```

Reporting

```
# Show perf.data in an ncurses browser:
perf report

# Show perf.data as a text report:
perf report --stdio

# List all events from perf.data:
perf script

# Annotate assembly instructions from perf.data
# with percentages
perf annotate [--stdio]
```

perf trace: trace system calls & other events

```
# Trace syscalls system-wide
perf trace

# Trace syscalls for PID
perf trace -p PID
```

perf record: record profiling data

records into
perf.data file

```
# Sample CPU functions for COMMAND, at 99 Hertz:
perf record -F 99 COMMAND

# Sample CPU functions for PID, until Ctrl-C:
perf record -p PID

# Sample CPU functions for PID, for 10 seconds:
perf record -p PID sleep 10

# Sample CPU stack traces for PID, for 10 seconds:
perf record -p PID -g -- sleep 10

# Sample CPU stack traces for PID, using DWARF to unwind stack:
perf record -p PID --call-graph dwarf
```

perf record: record tracing data

records into
perf.data file

```
# Trace new processes, until Ctrl-C:
perf record -e sched:sched_process_exec -a

# Trace all context-switches, until Ctrl-C:
perf record -e context-switches -a

# Trace all context-switches with stack traces, for 10 seconds:
perf record -e context-switches -ag -- sleep 10

# Trace all page faults with stack traces, until Ctrl-C:
perf record -e page-faults -ag
```

adding new trace events

```
# Add a tracepoint for kernel function tcp_sendmsg():
perf probe 'tcp_sendmsg'

# Trace previously created probe:
perf record -e -a probe:tcp_sendmsg

# Add a tracepoint for myfunc() return, and include the retval as a string:
perf probe 'myfunc%return +0($retval):string'

# Trace previous probe when size > 0, and state is not TCP_ESTABLISHED(1):
perf record -e -a probe:tcp_sendmsg --filter 'size > 0 && skc_state != 1' -a

# Add a tracepoint for do_sys_open() with the filename as a string:
perf probe 'do_sys_open filename:string'
```

these need
kernel debuginfo



perf_event_open()

- https://www.man7.org/linux/man-pages/man2/perf_event_open.2.html

