


OPENMP

OpenMP™

Open Multi-Processing

- Multi-*threaded* parallelism

<https://www.openmp.org/>



Source – Materials are from the Tutorials and Talks on OpenMP, Tim Mattson (Intel), Charles Augustine, Ruud Van der Paas (SUN), Ian Stoica


MOTIVATION

Multicore CPUs are everywhere:

- Servers with over 100 cores today
- Even smartphone CPUs have 8 cores

Multithreading, natural programming model

- All processors share the **same memory**
- Threads in a process see **same address space**
- OS scheduler decides when/which threads to run.
 - Typically threads are interleaved for fairness.



BUT...

Multithreading is hard

Lots of expertise necessary

Deadlocks and race conditions can occur.

Non-deterministic behavior makes it hard to debug

Synchronization is necessary to assure order and correct results.

Parallelize the following code using threads:

```
for (i=0; i<n; i++) {  
    sum = sum + sqrt(data[i]);  
}
```

Why hard?

Need mutex to protect the accesses to sum

Different code for serial and parallel version

No built-in tuning (# of processors?)

OPENMP: MULTI-PROCESSING PROGRAMMING

OpenMP: An API for Writing Multithreaded Applications

- A set of compiler directives and library routines for parallel application programmers
- Greatly simplifies writing multi-threaded (MT) programs in Fortran, C and C++
- Standardizes last 20 years of SMP practice

- Standard since 1997: *Fortran, C, C++*
- Requires Native language compiler to support OpenMP. – *most do*
- For shared memory machines: *Limited by available memory*

OPENMP

A language extension with constructs for parallel programming:

Directives & **C**lauses

Environment Variables

Functions

Work-sharing, Critical sections, atomic access, barriers, private variables

Parallelization is orthogonal to functionality

If the compiler does not recognize OpenMP directives, the code remains functional (albeit single-threaded).

Industry standard: supported by Intel, Microsoft, IBM, HP

INTRO TO OPENMP (C/C++)

- Preprocessor directives tell the compiler what to do
- Always start with #
- You've already seen one:

```
#include <stdio.h>
```

```
#include <omp.h>
```

- OpenMP directives tell the compiler to add machine code for parallel execution of the following *block*

```
#pragma omp parallel
```

- ➔ “Consider/Run the next set of instructions in parallel”

HELLO WORLD IN OPENMP!

```
#include <omp.h> //<-- necessary header file for OpenMP API
#include <stdio.h>


int main(void) {

    printf("OpenMP running with %d threads\n", omp_get_max_threads());

#pragma omp parallel //← Work Sharing Directives
    {
        //Code here will be executed by all threads
        printf("Hello World from thread %d\n", omp_get_thread_num());
    }
    return 0;
}
```

OPENMP EXECUTION MODEL

Fork and Join: Master thread spawns *a team of threads* as needed



OPENMP MEMORY MODEL

Shared memory model

Threads communicate by accessing shared variables

The sharing is defined syntactically


Any variable that is seen by two or more threads is shared

Any variable that is seen by one thread only is private

Race conditions possible

Use synchronization to protect from conflicts

Change how data is stored to minimize the synchronization



OPENMP – KEY VARIABLES AND SUBROUTINES

```
$export OMP_NUM_THREADS=X
```

- Environment variable to define Number of threads.

```
int omp_get_max_threads ()
```

- Returns max possible (generally set by OMP_NUM_THREADS)

```
int omp_get_num_threads ()
```

- Returns number of threads in a current team

```
int omp_get_thread_num ()
```

- Returns thread id of calling thread. (0 and omp_get_num_threads()]

```
int omp_get_max_procs ()
```

- Returns the number of processors on the machine.

```
double omp_get_wtime ()
```

- Returns the current reference time for the thread @current instruction.

OPENMP – CONSTRUCTS

Master

- Only executed by the Master/Main thread

Single

- Executed by only one of the thread in a team

Section/sections

- Declare block of instructions for task parallelism

atomic: Execute instruction atomically

barrier: Synchronization point for team of threads

critical: Only one thread at a time

OPENMP – CLAUSES

Default

- Default policy for variable sharing

Shared/private

- variables can be shared or private

Firstprivate/lastprivate

- Declare block of instructions for task parallelism

Reduction

- Scatter and Gather data across multiple threads

OPENMP: WORK SHARING EXAMPLE – DATA DECOMPOSITION

Sequential code

```
for (int i=0; i<N; i++) {sum = sum + a[i] + b[i];}
```

```
for (int i=0; i<N; i++) { a[i]=b[i]+c[i]; }
```

Semi-manual
parallelization

```
#pragma omp parallel
```

```
{
```

```
    int id = omp_get_thread_num();
```

```
    int nt = omp_get_num_threads();
```

```
    int i_start = id*N/nt, i_end = (id+1)*N/nt;
```

```
    for (int i=i_start; i<i_end; i++) { a[i]=b[i]+c[i]; }
```

```
}
```

- Launch **nt** threads
- Each thread uses **id** and **nt** variables to operate on a different segment of arrays

Automatic
parallelization with
#pragma omp for

```
#pragma omp parallel
```

```
{
```

```
    #pragma omp for
```

```
        for (int i=0; i<N; i++) { a[i]=b[i]+c[i]; }
```

```
}
```

OPENMP: WORK SHARING EXAMPLE – FUNCTIONAL DECOMPOSITION

```
answer1 = long_computation_1();  
answer2 = long_computation_2();  
if (answer1 != answer2) { ... }
```

How to parallelize?

```
#pragma omp sections  
{  
    #pragma omp section  
    answer1 = long_computation_1();  
    #pragma omp section  
    answer2 = long_computation_2();  
}  
if (answer1 != answer2) { ... }
```

OPENMP: DATA ENVIRONMENT

Shared Memory programming model

Most variables (including locals) are shared by threads

```
{  
    int sum = 0;  
    #pragma omp parallel for  
    for (int i=0; i<N; i++) sum += i;  
}
```

Global variables are **shared**

Some variables can be private

Variables inside the statement block

Variables in the called functions

Variables can be explicitly declared as private

QUICK CHECK!

```
int x = 5;
#pragma omp parallel num_threads(5)
{
    x++;
}
printf("X = %d\n", x);
```

```
int x = -1;
#pragma omp parallel num_threads(8)
{
    sleep(1);
    //Get thread number
    x = omp_get_thread_num();
}
printf("The value of x = %d\n", x);
```

- The same thing is done by all threads
- All data is shared between all threads
- Value of x at end of loop depends on .. (?)
- This code is non-deterministic and will produce different results on different runs

OVERRIDING STORAGE ATTRIBUTES

shared:

A single variable is shared across all threads
Correctness of updates is developer responsibility.

private:

A copy of the variable is created for each thread
There is no connection between original variable
and private copies
Can achieve same using variables inside { }

firstprivate:

Same, but the initial value of the variable is
copied from the main copy

lastprivate:

Same, but the last value of the variable is **copied to** the main copy

```
int i;  
#pragma omp parallel for private(i)  
for (i=0; i<n; i++) { ... }
```

```
int idx=1;  
int x = 10;  
#pragma omp parallel for \  
    firstprivate(x) lastprivate(idx)  
for (i=0; i<n; i++) {  
    if (data[i] == x)  
        idx = i;  
}
```

OMP PARALLEL CLAUSES 1

```
#pragma omp parallel if (scalar_expression)
```

- Only execute in parallel if true
- Otherwise serial

```
#pragma omp parallel private (list)
```

- Data local to thread
- Values are **not guaranteed to be defined on exit** (even if defined before)
- No storage associated with original object
 - Use `firstprivate` and/or `lastprivate` clause to override

OMP PARALLEL CLAUSES 2

```
#pragma omp parallel firstprivate (list)
```

- Variables in list are private
- Initialized with the value the variable had *before* entering the construct

```
#pragma omp parallel for lastprivate (list)
```

- Only in for loops
- Variables in list are private
- The thread that executes the *sequentially last iteration* updates the value of the variables in the list

OMP PARALLEL CLAUSE 3

```
#pragma omp shared (list)
```

- Data is accessible by all threads in team
- All threads access same address space
- Improperly scoped variables are big source of OMP bugs
 - Shared when should be private
 - Race condition

```
#pragma omp default (shared | none)
```

- Tip: Safest is to use default(none) and declare by hand

SHARED AND PRIVATE VARIABLES

- Take home message:
 - Be careful with the scope of your variables
 - Results must be independent of thread count
 - Test & debug thoroughly!
- Important note about compilers
 - C (before C99) does not allow variables declared in for loop syntax
 - Compiler will make loop variables private
 - Still recommend explicit

C

```
#pragma omp parallel private(i)
for (i=0; i<N; i++) {
    b = a + i;
}
```

C++

```
#pragma omp parallel
for (int i=0; i<N; i++) {
    b = a + i;
}
```

Automatically private

BEWARE OF RACE CONDITIONS!

```
for (int i=0; i<N; i++) {sum = sum + i;}
```

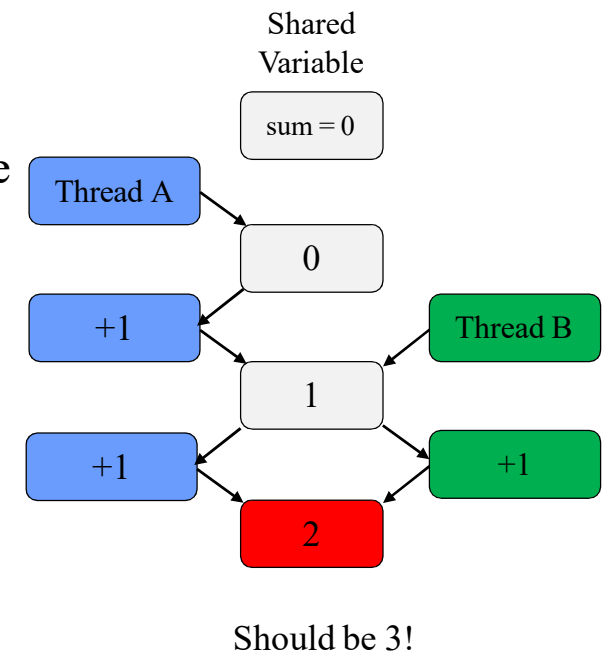
How to parallelize this code?
sum is not private!

When multiple threads simultaneously read/write shared variable

Multiple OMP solutions

- Reduction
- Atomic
- Critical

```
#pragma omp parallel for private(i) shared(sum)
for (i=0; i<N; i++) {
    sum += i;
}
```

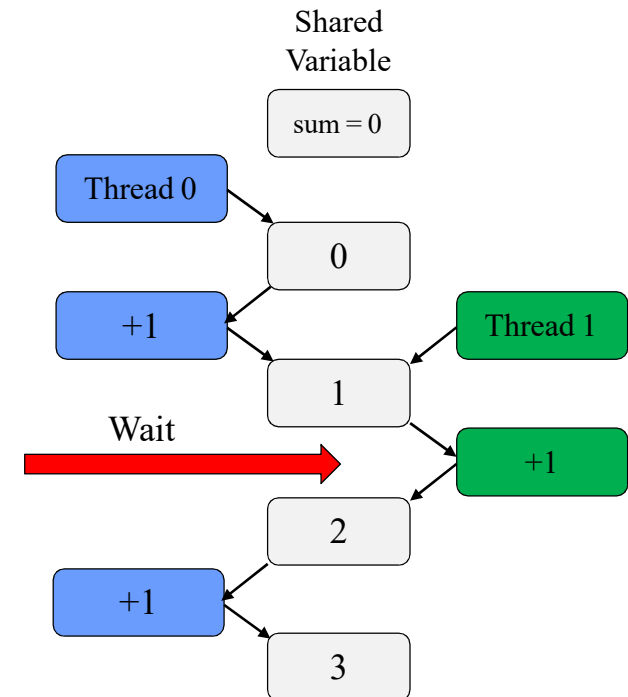


CRITICAL SECTION

- One solution: use critical
- Only one thread at a time can execute a critical section

```
#pragma omp critical
{
    sum += i;
}
```

- Downside?
 - SLOOOOOWWW
 - Overhead & serialization




OMP ATOMIC

- Atomic like “mini” critical
- Only one line
 - Certain limitations

```
#pragma omp atomic  
sum += i;
```

- Hardware controlled
 - Less overhead than critical

- Certain limitations...
 - Bonus point for the student(s) who identify the limitations



#PRAGMA OMP REDUCTION- REDUCTION

```
for (int i=0; i<N; i++) {sum = sum + a[i] * b[i];}
```

How to parallelize this code?

sum is not private, but accessing it atomically is too expensive

Have a private copy of sum in each thread, then add them up

Use the reduction clause

#pragma omp parallel for reduction(+: sum)

Any associative operator could be used: +, -, ||, |, *, etc

The private value is initialized automatically (to 0, 1, ...)

```
float dot_prod(float* a, float* b, int N)
{
    float sum = 0.0;
    #pragma omp parallel for reduction(+:sum)
    for(int i = 0; i < N; i++) {sum += a[i] * b[i]; }
    return sum;
}
```

Reduction:

Avoids race condition

Reduction variable must be shared

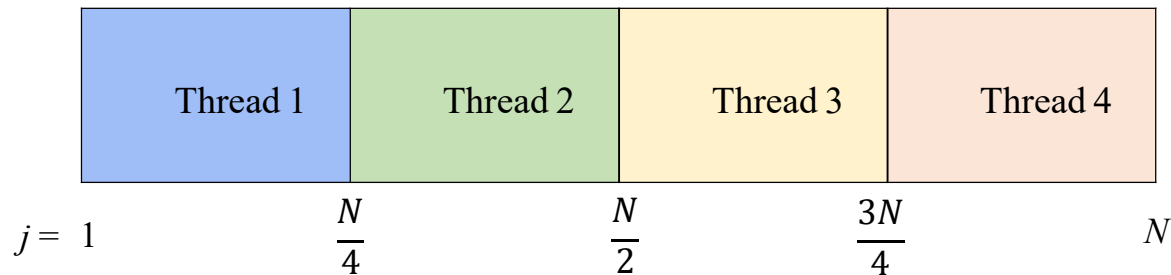
Makes variable private, then performs operator at end of loop

SCHEDULING OMP FOR

- How does a loop get split up?
 - In MPI, we have to do it manually.
- If you don't tell it what to do, the compiler decides
- Usually compiler chooses “static” – chunks of N/p

```
#pragma omp parallel for default(shared) private(j)
  for (j=0; j<N; j++) {
    ... // some work here
  }
```

Unspecified schedule




STATIC SCHEDULING

- You can tell the compiler what size chunks to take

```
#pragma omp parallel for default(shared)  
private(j) for (j=0; j<N; j++) {  
    ... // some work here  
}
```


`schedule(static,10)`



- Keeps assigning chunks until done
- Chunk size that isn't a multiple of the loop will result in threads with uneven numbers

PROBLEM WITH STATIC SCHEDULING

- What happens if loop iterations do not take the same amount of time?
 - Load imbalance




DYNAMIC SCHEDULING

- Chunks are assigned on the fly, as threads become available
 - When a thread finishes one chunk, it is assigned another

```
#pragma omp parallel for default(shared)  
private(j)  
for (j=0; j<N; j++) {  
    ... // some work here  
}
```

`schedule(dynamic,1)`



Note: It will have higher overhead than static!

OMP FOR SCHEDULING RECAP

```
#pragma omp parallel for schedule(type [,size])
```

- Scheduling types
 - Static
 - Chunks of specified size assigned round-robin
 - Dynamic
 - Chunks of specified size are assigned when thread finishes previous chunk
 - Guided
 - Like dynamic, but chunks are exponentially decreasing
 - Chunk will not be smaller than specified size
 - Runtime
 - Type and chunk determined at runtime via environment variables

SUMMARY

OpenMP: A framework for code parallelization

- Available for C/C++ and FORTRAN

- Based on a standard implementations from a wide selection of vendors

Relatively easy to use

- Write (and debug!) code first, parallelize later

- Parallelization can be incremental

- Parallelization can be turned off at runtime or compile time

- Code is still correct for a serial machine